

A Fixed Logic RISC Processor with Retargetable Infrastructure

Raul Garcia

The University of Manchester

Email: raul.garcia@manchester.ac.uk

Dirk Koch

The University of Manchester

Email: dirk.koch@manchester.ac.uk

Abstract—RISC processors implemented into fixed logic strive to deliver high performance for some compute-intensive applications. Typically, the solution is to extend the instruction set architecture (ISA) with specialized instructions. For instance, the ARM Cortex-A9 added the NEON SIMD engine for the media applications domain. This is a non-scalable solution in embedded systems with tight area and energy budgets. In this paper we propose an architecture where the area occupied by the fixed logic NEON engine is substituted by a retargetable infrastructure using FPGA logic. Arguably, our architecture has two key advantages. First, the parallel processing characteristics of the FPGA logic can be exploited by reconfiguring the infrastructure into a SIMD functional unit (FU). Our experimental results suggest that, for some media applications, a stripped-down version of a FPGA-based SIMD functional unit could match the performance of its fully-featured fixed logic counterpart. Second, the same infrastructure can be used to offload other compute-intensive functions from the processor into specialized FUs, targeted at other application domains. We investigated this idea and evaluated different software functions prone to become hardware-accelerated. We show resource utilization metrics of hardware-accelerated versions of this software functions. We argue that this collection of FUs could be implemented into the reconfigurable infrastructure within the limits of a strict FPGA fabric budget. This budget is estimated from the fixed logic NEON engine area.

I. INTRODUCTION

An Instruction Set Architecture (ISA) comprises all the operations that a processor can perform. Since the ISA of a RISC processor is designed to provide support for a wide variety of application domains, they are known as general purpose processors (GPP). Examples of the different application domains that a GPP can process are security, text processing, multimedia, networking, compression, etc. RISC processors are commonly implemented as hard processors in embedded systems. The term hard processor is used to describe a processor implemented into fixed logic (i.e. ASIC standard cells). As a result, the hardware of a hard processor can not be modified after the fabrication of the chip. For some applications, better performance is needed from RISC processors. This problem can be overcome by extending the ISA of the processor with an instruction or a set of instructions specially designed to boost the performance of that particular application. This concept is known as Instruction Set Extension (ISE). Each new ISE is application-driven. For instance, ARM extended the ISA of the ARMv6 with two application-specific instructions: USAD8 and USADA8. Respectively, they are used to calculate the unsigned sum of absolute differences and the accumulated unsigned sum of absolute differences between 8-bit values. These instructions are utilized for motion estimation, object recognition, and




Figure 1. SIMD Architecture. Same operation is performed on different elements, generating multiple data.

graphic compression algorithms such as MPEG. For the next iteration, the ARMv7, ARM introduced a Single Instruction Multiple Output (SIMD) ISE, known as the NEON engine. It was designed to accelerate multimedia and signal processing algorithms such as video encode/decode, 2D/3D graphics, gaming, audio and speech processing. Since these extensions are implemented into fixed logic, computer architects and hardware designers must release a new processor iteration each time a new ISE is implemented. Each new ISE has a cost in terms of chip area and energy consumption. Two disadvantages of a fixed logic ISE become evident. First, implementing a new ISE, impacts the time-to-market of the next product. Second, implementing new ISEs to accelerate more applications is a non-scalable solution, specially in embedded systems limited by strict area and energy consumption budgets. In this paper we present a different approach to deal with these disadvantages. We explore the idea of substituting the fixed logic area originally occupied by the NEON engine with fine-grained reconfigurable fabric. In theory, by making this substitution, we obtain a retargetable infrastructure within the hard processor.

A. The NEON SIMD Engine

The mobile market demands high performance from multimedia applications. In order to achieve this level of performance, most mobile processors feature a Single Instruction Multiple Output instruction set. A SIMD engine is a parallel architecture, where one single operation is executed by multiple processing units on multiple data elements. Figure 1 illustrates this architecture. ARM introduced on the ARMv7 architecture a dedicated SIMD processing unit known as NEON. The NEON SIMD engine has an independent pipeline and register file and provides 128-bit wide vector operations.

B. Reconfigurable Fabric

Field Programmable Gate Arrays (FPGAs) are semiconductor devices that are based around a matrix of Configurable Logic Blocks (CLBs). These CLBs are connected




Figure 5. The Xilinx Zynq Chip Architecture. The reconfigurable logic fabric is built around the fixed logic processor.

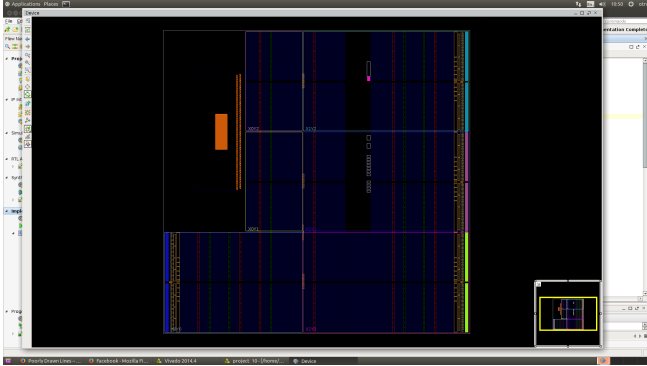


Figure 6. The Xilinx Zynq Chip Architecture. The reconfigurable logic fabric is built around the fixed logic processor.

same chip distributed in two sections: the Programmable Logic (PL), and the Processing system (PS).

2) *Budget estimation methodology*: Since the amount of building blocks (CLBs, Block RAM blocks, DSP slices) that the area corresponding to the PL section is known, it is possible to estimate the equivalence between the area that the PS represents and the amount of reconfigurable resources that could be allocated into that same area. For instance, using Xilinx tools to observe the floorplan of the XC7Z020 device, it is possible to count the number of CLB blocks implemented in Artix-7 FPGA fabric around the PS section.

3) *Results*: From the information obtained in the previous section, the height of the PS section in terms of CLB blocks can be estimated. From figure 6 it can be observed that the height of the PS section is about 50 CLBs (100 slices) high. Similarly, the PS width is about 16 CLB blocks (32 slices) wide, giving a total area of 800 CLBs. In the Xilinx 7 series FPGAs, four LUTs, their flip-flops as well as multiplexers and arithmetic carry logic form a slice, and two slices form a CLB [4], therefore the estimated amount of logic fabric that can be allocated into the PS area is about

6400 LUTs. Also, two columns of 20 BRAM blocks and two columns of 20 DSP slices must be considered. These estimations are summarized in table 1. Now that the amount of reconfigurable resources that the whole area of the PS section can allocate is known, it would be interesting to estimate how many resources the SIMD engine area only would represent. In order to make this estimation, we refer to Figure 1 and calculate the twenty percent of the resources previously estimated. The results represent approximately the amount of reconfigurable resources that can be allocated into the area currently dedicated to the SIMD engine in the ARM Cortex-A9. Since the NEON engine is part of the ARM Cortex-A9 hard IP, after the fabrication of the device, the hardware cannot be modified. However, if the hard IP area dedicated to the NEON engine were replaced with reconfigurable fabric, then this area would not be limited to perform media applications processing only, but it could be reconfigured into high-performance optimized engines for a wide range of application domains. This optimized engines would be bounded to the amount of resources that can be allocated into the same area, the estimation of this resources is presented in table 2. The implementation of this optimized engines would provide more flexibility and higher performance to the final product without adding area overhead.

FPGA Fabric	Artix-7 FPGA
CLBs	800
BRAM(36 Kb Blocks)	40
DSP slices	40

Table I
ARM PROCESSOR RESOURCE ESTIMATION

FPGA Fabric	Artix-7 FPGA
CLBs	160
BRAM(36 Kb Blocks)	8
DSP slices	8

Table II
SIMD ENGINE RESOURCE ESTIMATION

III. ARM PROCESSOR INSTRUCTIONS USAGE METRICS

We selected a set of representative software benchmarks from the multimedia application domain [2]. Also we selected other software benchmarks commonly used in embedded systems [3]. We used a GCC cross-compiler targeted at the ARMv7 architecture. We compiled with autovectorization enabled to generate SIMD instructions and executed the binaries on the Digilent Zybo board, that features a Xilinx Zynq chip with an ARM Cortex-A9. Then we developed a profiling framework based on the Valgrind dynamic analysis tool. Our primary objective was to study what vector instructions are being generated and which are more utilized. We obtained metrics for applications from multimedia and other application domains. The usage metrics are presented in Table III. From this information




Figure 7. The Xilinx Zynq Chip Architecture. The reconfigurable logic fabric is built around the fixed logic processor.

we can determine which instructions are good candidates to be implemented as functional units into the reconfigurable fabric.

Benchmark	Description	SIMD Instr. Executed
crc32	desc	NVINST1, VINST2
dijkstra	desc	VINST1, VINST2
fft	desc	VINST1, VINST2
lame	desc	VINST1, VINST2
rijndael	desc	VINST1, VINST2
sha	desc	VINST1, VINST2
rgn2gr	desc	VINST1, VINST2
rgn2lm	desc	VINST1, VINST2
vecadd	desc	VINST1, VINST2
poly	desc	VINST1, VINST2
matmul	desc	VINST1, VINST2
dotprod	desc	VINST1, VINST2
11lin	desc	VINST1, VINST2
11sat	desc	VINST1, VINST2

Table III
SIMD INSTRUCTION USAGE METRICS

IV. FPGA FABRIC RESOURCE UTILIZATION METRICS

In section 2 we estimated a strict resource budget based on the area of the fixed logic NEON engine. Now we will measure the amount of resources that different hardware benchmarks require in terms of reconfigurable fabric. This hardware benchmarks are selected from the study performed in the previous section. A stripped-down implementation of the NEON instruction set was developed. Our SIMD processing unit only supports integer operations and it uses the specialized "hard" DSP blocks of the FPGA fabric. We use this implementation as the baseline design to generate other metrics. The idea is that different, lighter versions of it can be used for specific functions heavily used in multimedia applications. A lighter implementation means less resource utilization and less energy consumption. Besides the hardware benchmarks targeted at media applications,

we also measured benchmarks for security, networking, and image processing. Table four shows that it is possible t

Benchmark	Description	FPGA resources utilized
crc32	desc	logic blocks
dijkstra	desc	logic blocks
fft	desc	logic blocks
lame	desc	logic blocks
rijndael	desc	logic blocks
sha	desc	logic blocks
rgn2gr	desc	logic blocks
rgn2lm	desc	logic blocks
vecadd	desc	logic blocks
poly	desc	logic blocks
matmul	desc	logic blocks
dotprod	desc	logic blocks
11lin	desc	logic blocks
11sat	desc	logic blocks

Table IV
FPGA FABRIC RESOURCE UTILIZATION METRICS

V. RELATED WORK

The idea of a system that combines the area efficiency and the performance of a Application-Specific Integrated Circuit (ASIC) with the reconfigurability of FPGAs has been around for some time. The GARP architecture [4] is one of the earliest works of this type. This architecture features a hard MIPS processors interconnected with a reconfigurable array. This array is dedicated to accelerate certain functions. Similarly, the One-Chip architecture [5] presents a system comprised of a RISC processor implemented into fixed logic and a loosely coupled array of FPGA fabric. This works show that it is possible to obtain benefits from coupling an efficient hard processor with reconfigurable functional units. However, this is not an easy task because, there exist a gap [6] between fixed logic and reconfigurable logic in terms of area, energy and speed. Later, as studied by [7], the idea of extending the ISA of RISC processors with custom instructions became popular. This custom instructions or instruction sets are implemented into reconfigurable fabric to hardware accelerate different software functions. The work of [8] and [9] show an interesting approach where a framework to detect hotspots and then automatically generate and implement a custom instruction. Frameworks to generate custom instructions based on commercial tools and popular RISC architectures such as the NIOS II soft processor [10] are part of a current trend that strive to close the gap between fixed logic and FPGA logic. In [11] the authors present the architecture of a framework for custom instruction generation. They build this architecture around the ARM processor [12]. One of the most popular and utilized architectures for embedded systems. They propose a microarchitectural interface to support a plug-and-play model for integrating a wide range of accelerator to a general purpose ARM processor. [13]

[14]

VI. CONCLUSIONS

In this paper we explored the idea of a RISC architecture where the fixed logic SIMD engine is substituted

by a FPGA-based reconfigurable infrastructure. We estimated a FPGA fabric budget for this infrastructure based on the current area of the fixed logic NEON engine. For its characteristics, the Xilinx Zynq chip was used as a base to make this estimation. Once we estimated a strict limit for the infrastructure budget, we selected a set of software benchmarks from different application domains to study what vector instructions are being generated by the GCC compiler. We also used the Valgrind profiling tool to expose software bottlenecks from those benchmarks. Then we presented resource utilization metrics from hardware functions to accelerate some of the software applications. We show that a collection of these functional units can be implemented into the reconfigurable fabric without exceeding the limits of the established budget. The future work will include a more detailed study of the integration of a reconfigurable functional unit to the fixed pipeline of the ARMv7 microarchitecture.

ACKNOWLEDGMENT

The research leading to these results has received funding from...

REFERENCES

- [1] L. H. Crockett, R. A. Elliot, M. A. Enderwitz, and R. W. Stewart, *The Zynq Book. Embedded Processing with the ARM Cortex-A9 on the Xilinx Zynq-7000 All Programmable SoC*. Strathclyde Academic Media, 2014.
- [2] C. Lee, M. Potkonjak, and W. H. Mangione-Smith, "MediaBench: A Tool for Evaluating and Synthesizing Multimedia and Communications Systems," *Thirtieth Annual IEEE/ACM International Symposium on Microarchitecture*, 1997., 2011.
- [3] M. R. Guthaus, T. M. A. Jeffrey S. Ringenberg, Dan Ernst, T. Mudge, and R. B. Brown, "MiBench: A free, commercially representative embedded benchmark suite," *2001 IEEE International Workshop on Workload Characterization*, 2001., 2011.
- [4] J. R. Hauser and J. Wawrzynek, "Garp: a MIPS processor with a reconfigurable coprocessor," *The 5th Annual IEEE Symposium on Field-Programmable Custom Computing Machines*, 1997.
- [5] R. D. Wittig and P. Chow, "OneChip: an FPGA processor with reconfigurable logic," *IEEE Symposium on FPGAs for Custom Computing Machines*, 1996.
- [6] I. Kuon and J. Rose, "Measuring the Gap Between FPGAs and ASICs," *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, 2012.
- [7] C. Galuzzi and K. Bertels, "The Instruction-Set Extension Problem: A Survey," *ACM Transactions on Reconfigurable Technology and Systems (TRETS)*, 2011.
- [8] A. Canis, J. Choi, M. Alham, V. Zhang, A. Kammoona, J. H. Anderson, S. Brown, and T. Czajkowski, "LegUp: high-level synthesis for FPGA-based processor/accelerator systems," *FPGA '11 Proceedings of the 19th ACM/SIGDA international symposium on Field programmable gate arrays*, 2011.
- [9] O. A. et al., "An End-to-End Design Flow for Automated Instruction Set Extension and Complex Instruction Selection Based on GCC," *GROW'09: Proceedings of the 1st International Workshop on GCC Research Opportunities*, 2009.
- [10] Altera Corp., "Nios II Custom Instruction. User Guide."
- [11] N. Clark, J. Blome, M. Chu, S. Mahlke, S. Biles, and K. Flautner, "An architecture framework for transparent instruction set customization in embedded processors," *32nd International Symposium on Computer Architecture, 2005. ISCA '05. Proceedings*, 2005.
- [12] ARM LTD., "Cortex-A9 Neon Media Processing Engine. Technical Reference Manual."
- [13] S. J. Jie and N. Kapre, "Comparing Soft and Hard Vector Processing in FPGA-based Embedded Systems," *2014 24th International Conference on Field Programmable Logic and Applications (FPL)*, 2014.
- [14] A. Severance and G. G. Lemieux, "Embedded supercomputing in FPGAs with the VectorBlox MXP Matrix Processor," in *2013 International Conference on Hardware/Software Code-sign and System Synthesis (CODES+ISSS)*, 2013.